

Elements Of Data Science - F2023

Week 12: Time Series

12/11/2023

TODOs

- Readings:
 - Recommended: DSFS: Chap 9: Getting Data
 - Recommended: DSFS: Chap 23: Databases and SQL
- HW4, **Due December 15 11:59pm ET**
- Final
 - Review sheet in github repo (soon!)
 - In class, in person.

Today

- Time Series Transformations

Questions?

Environment Setup

Environment Setup

```
In [1]: 1 import numpy
        2 import numpy as np
        3 import pandas as pd
        4 import matplotlib.pyplot as plt
        5 import seaborn as sns
        6
        7 sns.set_style('darkgrid')
        8 %matplotlib inline
```

Time Series

- Data ordered in time
- Applications
 - Financial
 - Economic
 - Scientific
 - etc.

Time Series Differences

- **Non-i.i.d.** : not independent and identically distributed
- not independent
 - Ex: Stock price
- not-identically distributed
 - Ex: Seasonality
- In other words: Order matters!

Representing Time in Python

- `datetime` library
- Pandas `Timestamp`


```
datetime.date
```

datetime.date

```
In [2]: 1 from datetime import date
        2
        3 friday = date(2022,11,1) # year,month,day
        4 friday
```

```
Out[2]: datetime.date(2022, 11, 1)
```

datetime.date

```
In [2]: 1 from datetime import date
        2
        3 friday = date(2022,11,1) # year,month,day
        4 friday
```

```
Out[2]: datetime.date(2022, 11, 1)
```

```
In [3]: 1 today = date.today()
        2 today
```

```
Out[3]: datetime.date(2023, 12, 10)
```

datetime.date

```
In [2]: 1 from datetime import date
        2
        3 friday = date(2022,11,1) # year,month,day
        4 friday
```

```
Out[2]: datetime.date(2022, 11, 1)
```

```
In [3]: 1 today = date.today()
        2 today
```

```
Out[3]: datetime.date(2023, 12, 10)
```

```
In [4]: 1 today.year
```

```
Out[4]: 2023
```

datetime.time

datetime.time

```
In [5]: 1 from datetime import time
        2
        3 class_start = time(19,10,0) # hour,minute,second,microsecond
        4 class_start
```

```
Out[5]: datetime.time(19, 10)
```

datetime.time

```
In [5]: 1 from datetime import time
        2
        3 class_start = time(19,10,0) # hour,minute,second,microsecond
        4 class_start
```

```
Out[5]: datetime.time(19, 10)
```

```
In [6]: 1 class_start.hour
```

```
Out[6]: 19
```

```
datetime.datetime
```


datetime.datetime

```
In [7]: 1 from datetime import datetime
        2
        3 # year,month,day,hour,minute,second,microsecond
        4 wednesday_afternoon = datetime(2022,11,30,15)
        5 wednesday_afternoon
```

```
Out[7]: datetime.datetime(2022, 11, 30, 15, 0)
```

datetime.datetime

```
In [7]: 1 from datetime import datetime
        2
        3 # year, month, day, hour, minute, second, microsecond
        4 wednesday_afternoon = datetime(2022, 11, 30, 15)
        5 wednesday_afternoon
```

```
Out[7]: datetime.datetime(2022, 11, 30, 15, 0)
```

```
In [8]: 1 now = datetime.now()
        2 now
```

```
Out[8]: datetime.datetime(2023, 12, 10, 22, 10, 40, 787053)
```

```
datetime.timedelta
```

datetime.timedelta

```
In [9]: 1 diff = datetime(2022,11,30,1) - datetime(2022,11,29,0)
        2 diff
```

```
Out[9]: datetime.timedelta(days=1, seconds=3600)
```

datetime.timedelta

```
In [9]: 1 diff = datetime(2022,11,30,1) - datetime(2022,11,29,0)
        2 diff
```

```
Out[9]: datetime.timedelta(days=1, seconds=3600)
```

```
In [10]: 1 diff.total_seconds()
```

```
Out[10]: 90000.0
```

datetime.timedelta

```
In [9]: 1 diff = datetime(2022,11,30,1) - datetime(2022,11,29,0)
        2 diff
```

```
Out[9]: datetime.timedelta(days=1, seconds=3600)
```

```
In [10]: 1 diff.total_seconds()
```

```
Out[10]: 90000.0
```

```
In [11]: 1 from datetime import timedelta
        2
        3 #days,seconds,microseconds,milliseconds,minutes,hours,weeks
        4 one_day = timedelta(1)
        5
        6 date(2022,11,30) + 2*one_day
```

```
Out[11]: datetime.date(2022, 12, 2)
```

Printing Datetimes: `strftime()`

Printing Datetimes: `strftime()`

```
In [12]: 1 now = datetime.now()  
        2 print(now)
```

```
2023-12-10 22:10:40.812133
```


Printing Datetimes: `strftime()`

```
In [12]: 1 now = datetime.now()  
        2 print(now)
```

```
2023-12-10 22:10:40.812133
```

```
In [13]: 1 now.strftime('%a %h %d, %Y %I:%M %p')
```

```
Out[13]: 'Sun Dec 10, 2023 10:10 PM'
```

Printing Datetimes: `strftime()`

```
In [12]: 1 now = datetime.now()  
        2 print(now)
```

```
2023-12-10 22:10:40.812133
```

```
In [13]: 1 now.strftime('%a %h %d, %Y %I:%M %p')
```

```
Out[13]: 'Sun Dec 10, 2023 10:10 PM'
```

```
1 %Y 4-digit year  
2 %y 2-digit year  
3 %m 2-digit month  
4 %d 2-digit day  
5 %H Hour (24-hour)  
6 %M 2-digit minute  
7 %S 2-digit second  
8 ...
```

Printing Datetimes: `strftime()`

```
In [12]: 1 now = datetime.now()  
        2 print(now)
```

```
2023-12-10 22:10:40.812133
```

```
In [13]: 1 now.strftime('%a %h %d, %Y %I:%M %p')
```

```
Out[13]: 'Sun Dec 10, 2023 10:10 PM'
```

```
1 %Y 4-digit year  
2 %y 2-digit year  
3 %m 2-digit month  
4 %d 2-digit day  
5 %H Hour (24-hour)  
6 %M 2-digit minute  
7 %S 2-digit second  
8 ...
```

See strftime.org and strfti.me

Parsing Datetimes: `pandas.to_datetime()`

- `dateutil.parser` available
- pandas has parser built in: `pd.to_datetime()`

Parsing Datetimes: `pandas.to_datetime()`

- `dateutil.parser` available
- pandas has parser built in: `pd.to_datetime()`

```
In [14]: 1 pd.to_datetime('11/30/2022 7:36pm')
```

```
Out[14]: Timestamp('2022-11-30 19:36:00')
```


pandas.Timestamp

- like `datetime.datetime`
- can include **timezone** and **frequency** info
- can handle a missing time: `NaT`
- can be used anywhere `datetime` can be used
- an array of Timestamps can be used as an index

pandas.Timestamp

- like `datetime.datetime`
- can include **timezone** and **frequency** info
- can handle a missing time: `NaT`
- can be used anywhere `datetime` can be used
- an array of Timestamps can be used as an index

```
In [16]: 1 pd.Timestamp(2022,11,30,19)
```

```
Out[16]: Timestamp('2022-11-30 19:00:00')
```


pandas.Timestamp

- like `datetime.datetime`
- can include **timezone** and **frequency** info
- can handle a missing time: `NaT`
- can be used anywhere `datetime` can be used
- an array of Timestamps can be used as an index

```
In [16]: 1 pd.Timestamp(2022,11,30,19)
```

```
Out[16]: Timestamp('2022-11-30 19:00:00')
```

```
In [17]: 1 pd.Timestamp('20221130 7:00pm EST')
```

```
Out[17]: Timestamp('2022-11-30 19:00:00-0500', tz='tzlocal()')
```

```
In [18]: 1 pd.Timestamp('20221130 7:00pm',tz='US/Pacific')
```

```
Out[18]: Timestamp('2022-11-30 19:00:00-0800', tz='US/Pacific')
```

pandas.Timestamp

- like `datetime.datetime`
- can include **timezone** and **frequency** info
- can handle a missing time: `NaT`
- can be used anywhere `datetime` can be used
- an array of Timestamps can be used as an index

```
In [16]: 1 pd.Timestamp(2022,11,30,19)
```

```
Out[16]: Timestamp('2022-11-30 19:00:00')
```

```
In [17]: 1 pd.Timestamp('20221130 7:00pm EST')
```

```
Out[17]: Timestamp('2022-11-30 19:00:00-0500', tz='tzlocal()')
```

```
In [18]: 1 pd.Timestamp('20221130 7:00pm',tz='US/Pacific')
```

```
Out[18]: Timestamp('2022-11-30 19:00:00-0800', tz='US/Pacific')
```

```
In [19]: 1 dt_index[0]
```

```
Out[19]: Timestamp('2020-11-26 00:00:00')
```

Accessing Datetime Components with `.dt`

Accessing Datetime Components with `.dt`

```
In [20]: 1 df_taxi = pd.read_csv('../data/yellowcab_tripdata_2017-01_subset10000rows.csv',  
2                                parse_dates=['tpep_pickup_datetime']).head(3)  
3 #df_taxi['tpep_pickup_datetime'] = pd.to_datetime(df_taxi.tpep_pickup_datetime)  
4 df_taxi.tpep_pickup_datetime
```

```
Out[20]: 0    2017-01-10 18:37:59  
1    2017-01-05 15:14:52  
2    2017-01-11 14:47:52  
Name: tpep_pickup_datetime, dtype: datetime64[ns]
```

Accessing Datetime Components with `.dt`

```
In [20]: 1 df_taxi = pd.read_csv('../data/yellowcab_tripdata_2017-01_subset10000rows.csv',
2                                     parse_dates=['tpep_pickup_datetime']).head(3)
3 #df_taxi['tpep_pickup_datetime'] = pd.to_datetime(df_taxi.tpep_pickup_datetime)
4 df_taxi.tpep_pickup_datetime
```

```
Out[20]: 0    2017-01-10 18:37:59
1    2017-01-05 15:14:52
2    2017-01-11 14:47:52
Name: tpep_pickup_datetime, dtype: datetime64[ns]
```

```
In [21]: 1 df_taxi.tpep_pickup_datetime.dt.day
```

```
Out[21]: 0     10
1       5
2     11
Name: tpep_pickup_datetime, dtype: int64
```

Accessing Datetime Components with `.dt`

```
In [20]: 1 df_taxi = pd.read_csv('../data/yellowcab_tripdata_2017-01_subset10000rows.csv',
2                               parse_dates=['tpep_pickup_datetime']).head(3)
3 #df_taxi['tpep_pickup_datetime'] = pd.to_datetime(df_taxi.tpep_pickup_datetime)
4 df_taxi.tpep_pickup_datetime
```

```
Out[20]: 0    2017-01-10 18:37:59
1    2017-01-05 15:14:52
2    2017-01-11 14:47:52
Name: tpep_pickup_datetime, dtype: datetime64[ns]
```

```
In [21]: 1 df_taxi.tpep_pickup_datetime.dt.day
```

```
Out[21]: 0     10
1       5
2     11
Name: tpep_pickup_datetime, dtype: int64
```

```
In [22]: 1 df_taxi.tpep_pickup_datetime.dt.day_of_week # Monday=0 ... Sunday=6
```

```
Out[22]: 0     1
1     3
2     2
Name: tpep_pickup_datetime, dtype: int64
```

Accessing Datetime Components with `.dt`

```
In [20]: 1 df_taxi = pd.read_csv('../data/yellowcab_tripdata_2017-01_subset10000rows.csv',
2                                     parse_dates=['tpep_pickup_datetime']).head(3)
3 #df_taxi['tpep_pickup_datetime'] = pd.to_datetime(df_taxi.tpep_pickup_datetime)
4 df_taxi.tpep_pickup_datetime
```

```
Out[20]: 0    2017-01-10 18:37:59
1    2017-01-05 15:14:52
2    2017-01-11 14:47:52
Name: tpep_pickup_datetime, dtype: datetime64[ns]
```

```
In [21]: 1 df_taxi.tpep_pickup_datetime.dt.day
```

```
Out[21]: 0     10
1       5
2     11
Name: tpep_pickup_datetime, dtype: int64
```

```
In [22]: 1 df_taxi.tpep_pickup_datetime.dt.day_of_week # Monday=0 ... Sunday=6
```

```
Out[22]: 0     1
1     3
2     2
Name: tpep_pickup_datetime, dtype: int64
```

```
In [23]: 1 df_taxi.tpep_pickup_datetime.dt.hour
```

```
Out[23]: 0     18
1     15
2     14
Name: tpep_pickup_datetime, dtype: int64
```

DateIndex Indexing/Selecting/Slicing

DateIndex Indexing/Selecting/Slicing

```
In [24]: 1 s = pd.Series(['Dec 1 2021', 'Jan 2 2022', 'Feb 3 2022'],  
2                 index=pd.to_datetime(['Dec 1 2021', 'Jan 2 2022', 'Feb 3 2022']))  
3 s
```

```
Out[24]: 2021-12-01    Dec 1 2021  
2022-01-02    Jan 2 2022  
2022-02-03    Feb 3 2022  
dtype: object
```

DateIndex Indexing/Selecting/Slicing

```
In [24]: 1 s = pd.Series(['Dec 1 2021', 'Jan 2 2022', 'Feb 3 2022'],  
2                     index=pd.to_datetime(['Dec 1 2021', 'Jan 2 2022', 'Feb 3 2022']))  
3 s
```

```
Out[24]: 2021-12-01    Dec 1 2021  
2022-01-02    Jan 2 2022  
2022-02-03    Feb 3 2022  
dtype: object
```

```
In [25]: 1 # can index normally using iloc  
2 s.iloc[0:2]
```

```
Out[25]: 2021-12-01    Dec 1 2021  
2022-01-02    Jan 2 2022  
dtype: object
```

DateIndex Indexing/Selecting/Slicing Cont.

DateIndex Indexing/Selecting/Slicing Cont.

```
In [26]: 1 # only rows from the year 2022  
        2 s.loc['2022']
```

```
Out[26]: 2022-01-02    Jan 2 2022  
        2022-02-03    Feb 3 2022  
        dtype: object
```

DateIndex Indexing/Selecting/Slicing Cont.

```
In [26]: 1 # only rows from the year 2022  
        2 s.loc['2022']
```

```
Out[26]: 2022-01-02    Jan 2 2022  
        2022-02-03    Feb 3 2022  
        dtype: object
```

```
In [27]: 1 # only rows from January 2022  
        2 s.loc['2022-01']
```

```
Out[27]: 2022-01-02    Jan 2 2022  
        dtype: object
```

DateIndex Indexing/Selecting/Slicing Cont.

```
In [26]: 1 # only rows from the year 2022  
        2 s.loc['2022']
```

```
Out[26]: 2022-01-02    Jan 2 2022  
        2022-02-03    Feb 3 2022  
        dtype: object
```

```
In [27]: 1 # only rows from January 2022  
        2 s.loc['2022-01']
```

```
Out[27]: 2022-01-02    Jan 2 2022  
        dtype: object
```

```
In [28]: 1 # only rows between Jan 1st 2021 and Jan 2nd 2022, inclusive  
        2 s.loc['01/01/2021':'01/02/2022']
```

```
Out[28]: 2021-12-01    Dec 1 2021  
        2022-01-02    Jan 2 2022  
        dtype: object
```

Datetimes in DataFrames

Datetimes in DataFrames

```
In [29]: 1 df = pd.DataFrame([[ '12/1/2021',101,'A'],
2                               [ '1/1/2022',102,'B']],columns=[ 'col1', 'col2', 'col3'])
3 df[ 'col1'] = pd.to_datetime(df.col1)
4 df.set_index( 'col1',drop=True,inplace=True)
5 df
```

Out[29]:

	col2	col3
col1		
2021-12-01	101	A
2022-01-01	102	B

Datetimes in DataFrames

```
In [29]: 1 df = pd.DataFrame([[ '12/1/2021',101,'A'],
2                        [ '1/1/2022',102,'B']],columns=[ 'col1', 'col2', 'col3'])
3 df[ 'col1'] = pd.to_datetime(df.col1)
4 df.set_index( 'col1',drop=True,inplace=True)
5 df
```

Out[29]:

	col2	col3
col1		
2021-12-01	101	A
2022-01-01	102	B

```
In [30]: 1 # only return rows from 2022
2 df.loc[ '2022']
```

Out[30]:

	col2	col3
col1		
2022-01-01	102	B

Timestamp Index: Setting Frequency

Timestamp Index: Setting Frequency

```
In [31]: 1 s = pd.Series(['Nov 1 2022', 'Nov 3 2022'], index=pd.to_datetime(['Nov 1 2022', 'Nov 3 2022']))  
        2 s
```

```
Out[31]: 2022-11-01    Nov 1 2022  
        2022-11-03    Nov 3 2022  
        dtype: object
```

Timestamp Index: Setting Frequency

```
In [31]: 1 s = pd.Series(['Nov 1 2022', 'Nov 3 2022'], index=pd.to_datetime(['Nov 1 2022', 'Nov 3 2022']))  
        2 s
```

```
Out[31]: 2022-11-01    Nov 1 2022  
        2022-11-03    Nov 3 2022  
        dtype: object
```

```
In [32]: 1 # Use resample() and asfreq() to set frequency  
        2 s.resample('D').asfreq()
```

```
Out[32]: 2022-11-01    Nov 1 2022  
        2022-11-02             NaN  
        2022-11-03    Nov 3 2022  
        Freq: D, dtype: object
```

Timestamp Index: Setting Frequency

```
In [31]: 1 s = pd.Series(['Nov 1 2022', 'Nov 3 2022'], index=pd.to_datetime(['Nov 1 2022', 'Nov 3 2022']))
          2 s
```

```
Out[31]: 2022-11-01    Nov 1 2022
          2022-11-03    Nov 3 2022
          dtype: object
```

```
In [32]: 1 # Use resample() and asfreq() to set frequency
          2 s.resample('D').asfreq()
```

```
Out[32]: 2022-11-01    Nov 1 2022
          2022-11-02             NaN
          2022-11-03    Nov 3 2022
          Freq: D, dtype: object
```

```
In [33]: 1 pd.to_datetime(['Nov 1 2022', 'Nov 3 2022'])
```

```
Out[33]: DatetimeIndex(['2022-11-01', '2022-11-03'], dtype='datetime64[ns]', freq=None)
```

Timestamp Index: Setting Frequency

```
In [31]: 1 s = pd.Series(['Nov 1 2022', 'Nov 3 2022'], index=pd.to_datetime(['Nov 1 2022', 'Nov 3 2022']))
          2 s
```

```
Out[31]: 2022-11-01    Nov 1 2022
          2022-11-03    Nov 3 2022
          dtype: object
```

```
In [32]: 1 # Use resample() and asfreq() to set frequency
          2 s.resample('D').asfreq()
```

```
Out[32]: 2022-11-01    Nov 1 2022
          2022-11-02             NaN
          2022-11-03    Nov 3 2022
          Freq: D, dtype: object
```

```
In [33]: 1 pd.to_datetime(['Nov 1 2022', 'Nov 3 2022'])
```

```
Out[33]: DatetimeIndex(['2022-11-01', '2022-11-03'], dtype='datetime64[ns]', freq=None)
```

```
In [34]: 1 # Use date_range with freq to get a range of dates of a certain frequency
          2 pd.date_range(start='Nov 1 2022', end='Nov 3 2022', freq='D')
```

```
Out[34]: DatetimeIndex(['2022-11-01', '2022-11-02', '2022-11-03'], dtype='datetime64[ns]', freq='D')
```

1	Sample of Available Frequencies		
2	B	business day frequency	
3	D	calendar day frequency	
4	W	weekly frequency	
5	M	month end frequency	
6	BM	business month end frequency	
7	...		
8	Q	quarter end frequency	
9	BQ	business quarter end frequency	
10	...		
11	Y	year end frequency	
12	BY	business year end frequency	
13	...		
14	BH	business hour frequency	
15	H	hourly frequency	
16	T,min	minutely frequency	
17	S	secondly frequency	
18	L,ms	milliseconds	
19	U,us	microseconds	
20	N	nanoseconds	

1	Sample of Available Frequencies		
2	B	business day frequency	
3	D	calendar day frequency	
4	W	weekly frequency	
5	M	month end frequency	
6	BM	business month end frequency	
7	...		
8	Q	quarter end frequency	
9	BQ	business quarter end frequency	
10	...		
11	Y	year end frequency	
12	BY	business year end frequency	
13	...		
14	BH	business hour frequency	
15	H	hourly frequency	
16	T,min	minutely frequency	
17	S	secondly frequency	
18	L,ms	milliseconds	
19	U,us	microseconds	
20	N	nanoseconds	

Timezones

- Handled by `pytz` library

Timezones

- Handled by `pytz` library

```
In [35]: 1 import pytz
          2
          3 [x for x in pytz.common_timezones if x.startswith('U')]
```

```
Out[35]: ['US/Alaska',
          'US/Arizona',
          'US/Central',
          'US/Eastern',
          'US/Hawaii',
          'US/Mountain',
          'US/Pacific',
          'UTC']
```

Timezones

- Handled by `pytz` library

```
In [35]: 1 import pytz
          2
          3 [x for x in pytz.common_timezones if x.startswith('U')]
```

```
Out[35]: ['US/Alaska',
          'US/Arizona',
          'US/Central',
          'US/Eastern',
          'US/Hawaii',
          'US/Mountain',
          'US/Pacific',
          'UTC']
```

UTC: coordinated universal time (EST is 5 hours behind, -5:00)

Timezones Cont.

Timezones Cont.

```
In [36]: 1 ts = pd.date_range('11/2/2022 9:30am', periods=2, freq='D')
          2 ts
```

```
Out[36]: DatetimeIndex(['2022-11-02 09:30:00', '2022-11-03 09:30:00'], dtype='datetime64[ns]', freq='D')
```

Timezones Cont.

```
In [36]: 1 ts = pd.date_range('11/2/2022 9:30am', periods=2, freq='D')
          2 ts
```

```
Out[36]: DatetimeIndex(['2022-11-02 09:30:00', '2022-11-03 09:30:00'], dtype='datetime64[ns]', freq='D')
```

```
In [37]: 1 # Set timezone using .tz_localize()
          2 ts_est = ts.tz_localize('US/Eastern')
          3 ts_est
```

```
Out[37]: DatetimeIndex(['2022-11-02 09:30:00-04:00', '2022-11-03 09:30:00-04:00'], dtype='datetime64[ns, US/Eastern]', freq=None)
```

Timezones Cont.

```
In [36]: 1 ts = pd.date_range('11/2/2022 9:30am', periods=2, freq='D')
          2 ts
```

```
Out[36]: DatetimeIndex(['2022-11-02 09:30:00', '2022-11-03 09:30:00'], dtype='datetime64[ns]', freq='D')
```

```
In [37]: 1 # Set timezone using .tz_localize()
          2 ts_est = ts.tz_localize('US/Eastern')
          3 ts_est
```

```
Out[37]: DatetimeIndex(['2022-11-02 09:30:00-04:00', '2022-11-03 09:30:00-04:00'], dtype='datetime64[ns, US/Eastern]', freq=None)
```

```
In [38]: 1 # Change timezones using .tz_convert()
          2 ts_est.tz_convert('UTC')
```

```
Out[38]: DatetimeIndex(['2022-11-02 13:30:00+00:00', '2022-11-03 13:30:00+00:00'], dtype='datetime64[ns, UTC]', freq=None)
```

Timezones Cont.

```
In [36]: 1 ts = pd.date_range('11/2/2022 9:30am', periods=2, freq='D')
        2 ts
```

```
Out[36]: DatetimeIndex(['2022-11-02 09:30:00', '2022-11-03 09:30:00'], dtype='datetime64[ns]', freq='D')
```

```
In [37]: 1 # Set timezone using .tz_localize()
        2 ts_est = ts.tz_localize('US/Eastern')
        3 ts_est
```

```
Out[37]: DatetimeIndex(['2022-11-02 09:30:00-04:00', '2022-11-03 09:30:00-04:00'], dtype='datetime64[ns, US/Eastern]', freq=None)
```

```
In [38]: 1 # Change timezones using .tz_convert()
        2 ts_est.tz_convert('UTC')
```

```
Out[38]: DatetimeIndex(['2022-11-02 13:30:00+00:00', '2022-11-03 13:30:00+00:00'], dtype='datetime64[ns, UTC]', freq=None)
```

```
In [39]: 1 # Can also initialize with timezone set
        2 ts = pd.date_range('11/2/2022 9:30am', periods=2, freq='D', tz='US/Eastern')
        3 ts
```

```
Out[39]: DatetimeIndex(['2022-11-02 09:30:00-04:00', '2022-11-03 09:30:00-04:00'], dtype='datetime64[ns, US/Eastern]', freq='D')
```


Time Series in Python so far:

- `datetime .date .time .datetime .timedelta`
- format with `.strftime()`
- parse time with `pd.to_datetime()`
- `pandas Timestamp Timedelta DatetimeIndex`
- Indexing with `DatetimeIndex`
- Frequencies
- Timezones

Next: Operations on Time Series data

- Shifting
- Resampling
- Moving Windows

Shifting/Lagging

- Moving data backward or forward in time (lagging/leading)
- Ex: calculate percent change

Shifting/Lagging

- Moving data backward or forward in time (lagging/leading)
- Ex: calculate percent change

```
In [40]: 1 ts = pd.Series([1,2,8],  
2                     index=pd.date_range('1/1/2022',periods=3,freq='M')) # Month End frequency (MS: Month Start)  
3 ts
```

```
Out[40]: 2022-01-31    1  
2022-02-28    2  
2022-03-31    8  
Freq: M, dtype: int64
```

Shifting/Lagging

- Moving data backward or forward in time (lagging/leading)
- Ex: calculate percent change

```
In [40]: 1 ts = pd.Series([1,2,8],  
2                     index=pd.date_range('1/1/2022',periods=3,freq='M')) # Month End frequency (MS: Month Start)  
3 ts
```

```
Out[40]: 2022-01-31    1  
2022-02-28    2  
2022-03-31    8  
Freq: M, dtype: int64
```

```
In [41]: 1 ts.shift(1) # last month's value
```

```
Out[41]: 2022-01-31    NaN  
2022-02-28    1.0  
2022-03-31    2.0  
Freq: M, dtype: float64
```

Shifting

- percent change, use one of :
 - $(\text{new_value} - \text{old_value}) / \text{old_value}$
 - $(\text{new_value} / \text{old_value}) - 1$

Shifting

- percent change, use one of :
 - $(\text{new_value} - \text{old_value}) / \text{old_value}$
 - $(\text{new_value} / \text{old_value}) - 1$

```
In [42]: 1 # multiply by 100 to turn into a percent
          2 ((ts / ts.shift(1)) - 1) * 100
```

```
Out[42]: 2022-01-31      NaN
          2022-02-28     100.0
          2022-03-31     300.0
          Freq: M, dtype: float64
```

Example Dataset: Twitter Stock

Example Dataset: Twitter Stock

```
In [43]: 1 # from pandas_datareader import data
2 # df_twtr = data.DataReader('TWTR', start='2015', end='11/27/2022', data_source='yahoo')
3 # df_twtr.to_csv('../data/twtr_20150102-20221127.csv')
4 df_twtr = pd.read_csv('../data/twtr_20150102-20221127.csv', parse_dates=['Date'], index_col='Date')
5 df_twtr.head(3).round(2)
```

Out[43]:

	High	Low	Open	Close	Volume	Adj Close
Date						
2015-01-02	36.74	35.54	36.23	36.56	12062461.0	36.56
2015-01-05	37.11	35.64	36.26	36.38	15062744.0	36.38
2015-01-06	39.45	36.04	36.27	38.76	33050812.0	38.76

Example Dataset: Twitter Stock

```
In [43]: 1 # from pandas_datareader import data
2 # df_twtr = data.DataReader('TWTR', start='2015', end='11/27/2022', data_source='yahoo')
3 # df_twtr.to_csv('../data/twtr_20150102-20221127.csv')
4 df_twtr = pd.read_csv('../data/twtr_20150102-20221127.csv', parse_dates=['Date'], index_col='Date')
5 df_twtr.head(3).round(2)
```

Out[43]:

	High	Low	Open	Close	Volume	Adj Close
Date						
2015-01-02	36.74	35.54	36.23	36.56	12062461.0	36.56
2015-01-05	37.11	35.64	36.26	36.38	15062744.0	36.38
2015-01-06	39.45	36.04	36.27	38.76	33050812.0	38.76


```
In [44]: 1 df_twtr.info() # Adj Close factors in corporate actions, such as stock splits, dividends, and rights offerings
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 1970 entries, 2015-01-02 to 2022-10-27
Data columns (total 6 columns):
#   Column      Non-Null Count  Dtype
---  -
0   High        1970 non-null   float64
1   Low         1970 non-null   float64
2   Open        1970 non-null   float64
3   Close       1970 non-null   float64
4   Volume      1970 non-null   float64
5   Adj Close   1970 non-null   float64
dtypes: float64(6)
memory usage: 107.7 KB
```

Example Dataset: Twitter Stock

Example Dataset: Twitter Stock

```
In [45]: 1 fig,ax = plt.subplots(1,1,figsize=(24,10))  
2 df_twtr[['Close']].plot(ax=ax,marker='x');  
3 ax.set_ylabel('stock price');
```



Shifting Example: Percent Change Twitter Close

Shifting Example: Percent Change Twitter Close

```
In [46]: 1 ((df_twtr.Close / df_twtr.Close.shift(1)) - 1).tail(3).round(3) # # (today / yesterday) - 1
```


```
Out[46]: Date
2022-10-25    0.024
2022-10-26    0.011
2022-10-27    0.007
Name: Close, dtype: float64
```

Shifting Example: Percent Change Twitter Close

```
In [46]: 1 ((df_twtr.Close / df_twtr.Close.shift(1)) - 1).tail(3).round(3) # # (today / yesterday) - 1
```

```
Out[46]: Date
2022-10-25    0.024
2022-10-26    0.011
2022-10-27    0.007
Name: Close, dtype: float64
```

```
In [47]: 1 # plot percent change of close in 2022
2 fig,ax = plt.subplots(1,1,figsize=(24,8))
3 close_2020 = df_twtr.loc['2020','Close']
4 ((close_2020 / close_2020.shift(1)) - 1).plot(marker='x',ax=ax,zorder=2);
5 ax.axhline(ls=':',c='k',zorder=1)
6 ax.set_ylabel('percent change');
```



Resampling

- Convert from one frequency to another
- **Downsampling**
 - from higher to lower (day to month)
 - need to aggregate
- **Upsampling**
 - from lower to higher (month to day)
 - need to fill missing
- Can also be used to set frequency from None

Resampling: Initialize Frequency

Resampling: Initialize Frequency

```
In [48]: 1 df_twtr.index
```

```
Out[48]: DatetimeIndex(['2015-01-02', '2015-01-05', '2015-01-06', '2015-01-07',  
                        '2015-01-08', '2015-01-09', '2015-01-12', '2015-01-13',  
                        '2015-01-14', '2015-01-15',  
                        ...  
                        '2022-10-14', '2022-10-17', '2022-10-18', '2022-10-19',  
                        '2022-10-20', '2022-10-21', '2022-10-24', '2022-10-25',  
                        '2022-10-26', '2022-10-27'],  
dtype='datetime64[ns]', name='Date', length=1970, freq=None)
```

Resampling: Initialize Frequency

```
In [48]: 1 df_twtr.index
```

```
Out[48]: DatetimeIndex(['2015-01-02', '2015-01-05', '2015-01-06', '2015-01-07',  
                        '2015-01-08', '2015-01-09', '2015-01-12', '2015-01-13',  
                        '2015-01-14', '2015-01-15',  
                        ...  
                        '2022-10-14', '2022-10-17', '2022-10-18', '2022-10-19',  
                        '2022-10-20', '2022-10-21', '2022-10-24', '2022-10-25',  
                        '2022-10-26', '2022-10-27'],  
                      dtype='datetime64[ns]', name='Date', length=1970, freq=None)
```

```
In [49]: 1 df_twtr_B = df_twtr.resample('B').asfreq() # set frequency to business day  
        2 df_twtr_B.index
```

```
Out[49]: DatetimeIndex(['2015-01-02', '2015-01-05', '2015-01-06', '2015-01-07',  
                        '2015-01-08', '2015-01-09', '2015-01-12', '2015-01-13',  
                        '2015-01-14', '2015-01-15',  
                        ...  
                        '2022-10-14', '2022-10-17', '2022-10-18', '2022-10-19',  
                        '2022-10-20', '2022-10-21', '2022-10-24', '2022-10-25',  
                        '2022-10-26', '2022-10-27'],  
                      dtype='datetime64[ns]', name='Date', length=2040, freq='B')
```

Resampling: Downsampling

- Go from higher/shorter to lower/longer
- Need to aggregate (like groupby)
- Example: Downsampling from business day to business quarter

Resampling: Downsampling

- Go from higher/shorter to lower/longer
- Need to aggregate (like groupby)
- Example: Downsampling from business day to business quarter

```
In [50]: 1 df_twtr_BQ = df_twtr_B.resample('BQ')  
        2 df_twtr_BQ
```

```
Out[50]: <pandas.core.resample.DatetimeIndexResampler object at 0x7fdef51cc250>
```

Resampling: Downsampling

- Go from higher/shorter to lower/longer
- Need to aggregate (like groupby)
- Example: Downsampling from business day to business quarter

```
In [50]: 1 df_twtr_BQ = df_twtr_B.resample('BQ')
          2 df_twtr_BQ
```

```
Out[50]: <pandas.core.resample.DatetimeIndexResampler object at 0x7fdef51cc250>
```

```
In [51]: 1 print(df_twtr_BQ)
```

```
DatetimeIndexResampler [freq=<BusinessQuarterEnd: startingMonth=12>, axis=0, closed=right, label=right, convention=start, origin=start_day]
```

Resampling: Downsampling

- Go from higher/shorter to lower/longer
- Need to aggregate (like groupby)
- Example: Downsampling from business day to business quarter

```
In [50]: 1 df_twtr_BQ = df_twtr_B.resample('BQ')
        2 df_twtr_BQ
```

```
Out[50]: <pandas.core.resample.DatetimeIndexResampler object at 0x7fdef51cc250>
```

```
In [51]: 1 print(df_twtr_BQ)
```

```
DatetimeIndexResampler [freq=<BusinessQuarterEnd: startingMonth=12>, axis=0, closed=right, label=right, convention=start, origin=start_day]
```

```
In [52]: 1 df_twtr_BQ.mean().head(3).round(2)
```

```
Out[52]:
```

	High	Low	Open	Close	Volume	Adj Close
Date						
2015-03-31	45.10	43.55	44.23	44.34	20840997.51	44.34
2015-06-30	41.63	40.38	41.17	40.87	22287099.56	40.87
2015-09-30	30.64	29.42	30.05	30.00	20065038.11	30.00

Resampling: Downsampling

Resampling: Downsampling

```
In [53]: 1 fig,ax = plt.subplots(1,1,figsize=(24,8))
2 df_twtr_B.Close.plot(style='-', label='by B',ax=ax)
3 df_twtr_BQ.Close.mean().plot(style='--',marker='x',label='by BQ',ax=ax)
4 plt.legend(loc='upper right');
```



Resampling: Upsampling

- Go from lower/longer to higher/shorter
- Need to decide how to handle missing values
- Example: Upsample from business day to hour

Resampling: Upsampling

- Go from lower/longer to higher/shorter
- Need to decide how to handle missing values
- Example: Upsample from business day to hour

```
In [54]: 1 df_twtr_B.index[:3]
```

```
Out[54]: DatetimeIndex(['2015-01-02', '2015-01-05', '2015-01-06'], dtype='datetime64[ns]', name='Date', freq='B')
```

Resampling: Upsampling

- Go from lower/longer to higher/shorter
- Need to decide how to handle missing values
- Example: Upsample from business day to hour

```
In [54]: 1 df_twtr_B.index[:3]
```

```
Out[54]: DatetimeIndex(['2015-01-02', '2015-01-05', '2015-01-06'], dtype='datetime64[ns]', name='Date', freq='B')
```

```
In [55]: 1 df_twtr_B.Close.resample('H').asfreq().iloc[0:3].round(2)
```

```
Out[55]: Date
2015-01-02 00:00:00    36.56
2015-01-02 01:00:00     NaN
2015-01-02 02:00:00     NaN
Freq: H, Name: Close, dtype: float64
```

Resampling: Upsampling

- Go from lower/longer to higher/shorter
- Need to decide how to handle missing values
- Example: Upsample from business day to hour

```
In [54]: 1 df_twtr_B.index[:3]
```

```
Out[54]: DatetimeIndex(['2015-01-02', '2015-01-05', '2015-01-06'], dtype='datetime64[ns]', name='Date', freq='B')
```

```
In [55]: 1 df_twtr_B.Close.resample('H').asfreq().iloc[0:3].round(2)
```

```
Out[55]: Date
2015-01-02 00:00:00    36.56
2015-01-02 01:00:00     NaN
2015-01-02 02:00:00     NaN
Freq: H, Name: Close, dtype: float64
```

```
In [56]: 1 df_twtr_B.Close.resample('H').asfreq().iloc[70:73].round(2)
```

```
Out[56]: Date
2015-01-04 22:00:00     NaN
2015-01-04 23:00:00     NaN
2015-01-05 00:00:00    36.38
Freq: H, Name: Close, dtype: float64
```

Resampling: Upsampling

- `ffill()` : Forward Fill

Resampling: Upsampling

- `ffill()` : Forward Fill

```
In [57]: 1 df_twtr_B.Close.resample('H').ffill().head(3).round(2)
```

```
Out[57]: Date
2015-01-02 00:00:00    36.56
2015-01-02 01:00:00    36.56
2015-01-02 02:00:00    36.56
Freq: H, Name: Close, dtype: float64
```

Resampling: Upsampling

- `ffill()` : Forward Fill

```
In [57]: 1 df_twtr_B.Close.resample('H').ffill().head(3).round(2)
```

```
Out[57]: Date
2015-01-02 00:00:00    36.56
2015-01-02 01:00:00    36.56
2015-01-02 02:00:00    36.56
Freq: H, Name: Close, dtype: float64
```

- `bfill()` : Backward Fill

Resampling: Upsampling

- `ffill()` : Forward Fill

```
In [57]: 1 df_twtr_B.Close.resample('H').ffill().head(3).round(2)
```

```
Out[57]: Date
2015-01-02 00:00:00    36.56
2015-01-02 01:00:00    36.56
2015-01-02 02:00:00    36.56
Freq: H, Name: Close, dtype: float64
```

- `bfill()` : Backward Fill

```
In [58]: 1 df_twtr_B.Close.resample('H').bfill().head(3).round(3)
```

```
Out[58]: Date
2015-01-02 00:00:00    36.56
2015-01-02 01:00:00    36.38
2015-01-02 02:00:00    36.38
Freq: H, Name: Close, dtype: float64
```


Moving/Rolling Windows

- Apply function on a fixed window moving across time
- Method of smoothing out the data
- **center** : place values at center of window

Moving/Rolling Windows

- Apply function on a fixed window moving across time
- Method of smoothing out the data
- **center** : place values at center of window

```
In [59]: 1 df_twtr_B.Close['2020-11-02':'2020-11-06'].round(2)
```

```
Out[59]: Date
2020-11-02    39.47
2020-11-03    41.73
2020-11-04    42.76
2020-11-05    43.71
2020-11-06    43.12
Freq: B, Name: Close, dtype: float64
```

Moving/Rolling Windows

- Apply function on a fixed window moving across time
- Method of smoothing out the data
- **center** : place values at center of window

```
In [59]: 1 df_twtr_B.Close['2020-11-02':'2020-11-06'].round(2)
```

```
Out[59]: Date
2020-11-02    39.47
2020-11-03    41.73
2020-11-04    42.76
2020-11-05    43.71
2020-11-06    43.12
Freq: B, Name: Close, dtype: float64
```

```
In [60]: 1 rolling_3 = df_twtr_B.Close['2020-11-02':'2020-11-06'].rolling(3, center=True)
2 rolling_3
```

```
Out[60]: Rolling [window=3,center=True,axis=0,method=single]
```

Moving/Rolling Windows

- Apply function on a fixed window moving across time
- Method of smoothing out the data
- **center** : place values at center of window

```
In [59]: 1 df_twtr_B.Close['2020-11-02':'2020-11-06'].round(2)
```

```
Out[59]: Date
2020-11-02    39.47
2020-11-03    41.73
2020-11-04    42.76
2020-11-05    43.71
2020-11-06    43.12
Freq: B, Name: Close, dtype: float64
```

```
In [60]: 1 rolling_3 = df_twtr_B.Close['2020-11-02':'2020-11-06'].rolling(3, center=True)
2 rolling_3
```

```
Out[60]: Rolling [window=3,center=True,axis=0,method=single]
```


```
In [61]: 1 rolling_3.mean()['2020-11-02':'2020-11-06'].round(2)
```

```
Out[61]: Date
2020-11-02    NaN
2020-11-03    41.32
2020-11-04    42.73
2020-11-05    43.20
2020-11-06    NaN
Freq: B, Name: Close, dtype: float64
```

Moving Windows

Moving Windows

```
In [62]: 1 sns.set_style("whitegrid")
2         rolling = df_twtr_B.Close.rolling(5, center=True)
3
4         fig,ax = plt.subplots(1,1,figsize=(24,8));
5         df_twtr_B.loc['2020'].Close.plot(style='-',alpha=0.3,label='business day');
6         rolling.mean().loc['2020'].plot(style='--',label='5 day rolling window mean');
7         (rolling.mean().loc['2020'] + 2*rolling.std().loc['2020']).plot(style=':',c='g',label='_nolegend_');
8         (rolling.mean().loc['2020'] - 2*rolling.std().loc['2020']).plot(style=':',c='g',label='_nolegend_');
9         ax.legend();
```



Example: Bike Travel (From PDSH Chapter 3.11)

- Bicycle traffic over Fremont Bridge in Seattle in 2012
- Data gathered using: `!curl -o ../data/FremontBridge.csv https://data.seattle.gov/api/views/65db-xm6k/rows.csv?accessType=DOWNLOAD`

Example: Bike Travel (From PDSH Chapter 3.11)

- Bicycle traffic over Fremont Bridge in Seattle in 2012
- Data gathered using: `!curl -o ../data/FremontBridge.csv https://data.seattle.gov/api/views/65db-xm6k/rows.csv?accessType=DOWNLOAD`

```
In [63]: 1 df_bike_counts = pd.read_csv('../data/FremontBridge_2012-2015.csv', parse_dates=['Date'], index_col='Date')
          2 df_bike_counts.columns = ['Total', 'East', 'West']
          3 df_bike_counts.info()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 28440 entries, 2012-10-03 00:00:00 to 2015-12-31 23:00:00
Data columns (total 3 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Total    28433 non-null     float64
1   East     28433 non-null     float64
2   West     28433 non-null     float64
dtypes: float64(3)
memory usage: 888.8 KB
```


Example: Bike Travel (From PDSH Chapter 3.11)

- Bicycle traffic over Fremont Bridge in Seattle in 2012
- Data gathered using: `!curl -o ../data/FremontBridge.csv https://data.seattle.gov/api/views/65db-xm6k/rows.csv?accessType=DOWNLOAD`

```
In [63]: 1 df_bike_counts = pd.read_csv('../data/FremontBridge_2012-2015.csv', parse_dates=['Date'], index_col='Date')
          2 df_bike_counts.columns = ['Total', 'East', 'West']
          3 df_bike_counts.info()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 28440 entries, 2012-10-03 00:00:00 to 2015-12-31 23:00:00
Data columns (total 3 columns):
#   Column  Non-Null Count  Dtype  
---  -
0   Total    28433 non-null     float64
1   East     28433 non-null     float64
2   West     28433 non-null     float64
dtypes: float64(3)
memory usage: 888.8 KB
```

```
In [64]: 1 df_bike_counts.head(3)
```

Out[64]:

	Total	East	West
Date			
2012-10-03 00:00:00	13.0	4.0	9.0
2012-10-03 01:00:00	10.0	4.0	6.0
2012-10-03 02:00:00	2.0	1.0	1.0

Example: Fill Missing Values

Example: Fill Missing Values

```
In [65]: 1 f'proportion missing: {sum(df_bike_counts.Total.isna()) / len(df_bike_counts):0.5f}'
```

```
Out[65]: 'proportion missing: 0.00025'
```

Example: Fill Missing Values

```
In [65]: 1 f'proportion missing: {sum(df_bike_counts.Total.isna()) / len(df_bike_counts):0.5f}'
```

```
Out[65]: 'proportion missing: 0.00025'
```


```
In [66]: 1 df_bike_counts = df_bike_counts.fillna(method='ffill')
2 df_bike_counts.info()
```

```
<class 'pandas.core.frame.DataFrame'>
DatetimeIndex: 28440 entries, 2012-10-03 00:00:00 to 2015-12-31 23:00:00
Data columns (total 3 columns):
#   Column  Non-Null Count  Dtype
---  -
0   Total    28440 non-null    float64
1   East     28440 non-null    float64
2   West     28440 non-null    float64
dtypes: float64(3)
memory usage: 888.8 KB
```

Plot data from 2015

Plot data from 2015


```
In [67]: 1 fig,ax = plt.subplots(1,1,figsize=(24,8))
2 df_bike_counts.loc['2015'].plot(style=['-', '--', ':'],ax=ax)
3 plt.ylabel('Hourly Bicycle Count');
```



Downsample to weekly sum to smooth things out

Downsample to weekly sum to smooth things out


```
In [68]: 1 weekly = df_bike_counts.resample('W').sum()  
2 weekly.plot(style=[':', '--', '-'], figsize=(24,8))  
3 plt.ylabel('Weekly bicycle count');
```



Resample at daily for a more granular view and apply a rolling window of 30 days

Resample at daily for a more granular view and apply a rolling window of 30 days


```
In [69]: 1 daily = df_bike_counts.resample('D').sum()  
2 daily.rolling(30,center=True).mean().plot(style=[':', '--', '-'], figsize=(24,8))  
3 plt.ylabel('mean daily count');
```



A wider window using a gaussian filter smooths more while accentuating daily differences

A wider window using a gaussian filter smooths more while accentuating daily differences

```
In [70]: 1 daily.rolling(30,center=True,win_type='gaussian').mean(std=7).plot(style=[':', '--', '-'],figsize=(24,8));  
2 plt.ylabel('mean daily count');
```



From Datetime to Time

From Datetime to Time

```
In [71]: 1 #If we want to only look at time of day  
        2 df_bike_counts.index.time
```

```
Out[71]: array([datetime.time(0, 0), datetime.time(1, 0), datetime.time(2, 0), ...,  
               datetime.time(21, 0), datetime.time(22, 0), datetime.time(23, 0)],  
              dtype=object)
```

From Datetime to Time

```
In [71]: 1 #If we want to only look at time of day
         2 df_bike_counts.index.time
```

```
Out[71]: array([datetime.time(0, 0), datetime.time(1, 0), datetime.time(2, 0), ...,
               datetime.time(21, 0), datetime.time(22, 0), datetime.time(23, 0)],
              dtype=object)
```


```
In [72]: 1 #Get mean data by time (hourly)
         2 by_time = df_bike_counts.groupby(df_bike_counts.index.time).mean().round(2)
         3 display(by_time)
```

	Total	East	West
00:00:00	13.34	5.94	7.40
01:00:00	7.15	3.34	3.81
02:00:00	4.97	2.61	2.36
03:00:00	3.43	1.90	1.52
04:00:00	6.13	3.53	2.59
05:00:00	24.26	16.22	8.04
06:00:00	87.65	58.94	28.71
07:00:00	212.38	154.07	58.31
08:00:00	298.85	206.76	92.09
09:00:00	182.88	115.71	67.17
10:00:00	97.45	57.71	39.74
11:00:00	85.06	48.46	36.61
12:00:00	91.21	50.26	40.96
13:00:00	97.83	52.33	45.50
14:00:00	106.61	55.18	51.42
15:00:00	131.29	63.61	67.67
16:00:00	207.88	86.06	121.82
17:00:00	361.78	117.19	244.59
18:00:00	264.50	96.03	168.47
19:00:00	129.85	52.96	76.89

Plot by hour of the day

Plot by hour of the day


```
In [73]: 1 hourly_ticks = 60 * 60 * 4 * np.arange(6) # sec * min * every4hours
2 by_time.plot(xticks=hourly_ticks, style=[':', '--', '-'], figsize=(24,8));
3 plt.ylabel('mean hourly count');
```



Can also look at average by day of week

Can also look at average by day of week



```
In [74]: 1 # note that for dayofweek: 0 == Mon, 1 == Tues,..., 6 == 'Sun'
2 by_weekday = df_bike_counts.groupby(df_bike_counts.index.dayofweek).mean()
3 by_weekday = by_weekday.set_index(pd.Index(['Mon', 'Tues', 'Wed', 'Thurs', 'Fri', 'Sat', 'Sun']))
4
5 fig,ax = plt.subplots(1,1,figsize=(24,8))
6 by_weekday.plot(style=[':', '--', '-'], ax=ax);
7 ax.set_xlabel('Day of Week');ax.set_ylabel('mean daily count');
```



Separate out weekdays and weekends

Separate out weekdays and weekends

```
In [75]: 1 # create a weekend mask
2 weekend = np.where(df_bike_counts.index.day_of_week < 5, 'Weekday', 'Weekend')
3
4 # get hourly mean values split by weekday, weekend
5 by_time = df_bike_counts.groupby([weekend, df_bike_counts.index.time]).mean()
6 fig, ax = plt.subplots(1, 2, figsize=(24, 8))
7 by_time.loc['Weekday'].plot(ax=ax[0], title='Weekdays', xticks=hourly_ticks, style=[':', '--', '-'])
8 by_time.loc['Weekend'].plot(ax=ax[1], title='Weekends', xticks=hourly_ticks, style=[':', '--', '-']);
```



Can we predict daily Total bike traffic?

Can we predict daily Total bike traffic?

```
In [76]: 1 df_bike_counts = pd.read_csv('../data/FremontBridge_2012-2015.csv', index_col='Date', parse_dates=True)
2 df_bike_weather = pd.read_csv('../data/BicycleWeather.csv', index_col='DATE', parse_dates=True)
3
4 df_bike = (
5     df_bike_counts.loc[:, ['Fremont Bridge Total']] # keep Total as target
6     .rename({'Fremont Bridge Total': 'Total'}, axis=1) # rename target column
7     .resample('D').sum() # downsample to daily totals
8 )
9 print(df_bike.head(3))
```

Date	Total
2012-10-03	3521.0
2012-10-04	3475.0
2012-10-05	3148.0

Can we predict daily Total bike traffic?

```
In [76]: 1 df_bike_counts = pd.read_csv('../data/FremontBridge_2012-2015.csv', index_col='Date', parse_dates=True)
2 df_bike_weather = pd.read_csv('../data/BicycleWeather.csv', index_col='DATE', parse_dates=True)
3
4 df_bike = (
5     df_bike_counts.loc[:, ['Fremont Bridge Total']] # keep Total as target
6     .rename({'Fremont Bridge Total': 'Total'}, axis=1) # rename target column
7     .resample('D').sum() # downsample to daily totals
8 )
9 print(df_bike.head(3))
```

Date	Total
2012-10-03	3521.0
2012-10-04	3475.0
2012-10-05	3148.0

On to Feature Engineering...

Add 'day of week'

Add 'day of week'

```
In [77]: 1 day_names_map = dict(enumerate(['Mon', 'Tue', 'Wed', 'Thu', 'Fri', 'Sat', 'Sun']))
        2 print(f"{day_names_map = :}")
        3 df_bike['DayOfWeek'] = df_bike.index.dayofweek.map(day_names_map)
        4 df_bike.head(3)
```

day_names_map = {0: 'Mon', 1: 'Tue', 2: 'Wed', 3: 'Thu', 4: 'Fri', 5: 'Sat', 6: 'Sun'}

Out[77]:

	Total	DayOfWeek
Date		
2012-10-03	3521.0	Wed
2012-10-04	3475.0	Thu
2012-10-05	3148.0	Fri

Add 'is it a holiday' dummy feature

Add 'is it a holiday' dummy feature


```
In [78]: 1 from pandas.tseries.holiday import USFederalHolidayCalendar
2 cal = USFederalHolidayCalendar()
3 holidays = cal.holidays('2012', '2016')
4
5 df_bike = df_bike.join(pd.Series(1, index=holidays, name='IsHoliday'))
6 df_bike['IsHoliday'].fillna(0, inplace=True)
7 print(df_bike.head(3))
```

	Total	DayOfWeek	IsHoliday
Date			
2012-10-03	3521.0	Wed	0.0
2012-10-04	3475.0	Thu	0.0
2012-10-05	3148.0	Fri	0.0

Add number of hours of daylight

Add number of hours of daylight

```
In [79]: 1 from datetime import datetime
2
3 def hours_of_daylight(date, axis=23.44, latitude=47.61):
4     """Compute the hours of daylight for the given date"""
5     days = (date - datetime(2000, 12, 21)).days # days till winter solstice
6     m = (1. - np.tan(np.radians(latitude))
7          * np.tan(np.radians(axis) * np.cos(days * 2 * np.pi / 365.25)))
8     return 24. * np.degrees(np.arccos(1 - np.clip(m, 0, 2))) / 180.
9
10 df_bike['HoursDaylight'] = list(map(hours_of_daylight, df_bike.index));
11
12 ax = df_bike[['HoursDaylight']].plot(figsize=(18,4));
13 ax.set_ylim(8, 16);
```



Add weather information (Q: can we predict this for future dates?)

Add weather information (Q: can we predict this for future dates?)

```
In [80]: 1 # temperatures are in 1/10 deg C; convert to C
2 df_bike_weather['TMIN'] /= 10
3 df_bike_weather['TMAX'] /= 10
4 df_bike_weather['TempC'] = 0.5 * (df_bike_weather['TMIN'] + df_bike_weather['TMAX'])
5
6 # precip is in 1/10 mm; convert to inches
7 df_bike_weather['PRCP'] /= 254
8 df_bike_weather['IsDryDay'] = (df_bike_weather['PRCP'] == 0).astype(int)
9
10 df_bike = df_bike.join(df_bike_weather[['PRCP', 'TempC', 'IsDryDay']],how='inner')
11 df_bike.head(3).round(2)
```

Out[80]:

	Total	DayOfWeek	IsHoliday	HoursDaylight	PRCP	TempC	IsDryDay
2012-10-03	3521.0	Wed	0.0	11.28	0.0	13.35	1
2012-10-04	3475.0	Thu	0.0	11.22	0.0	13.60	1
2012-10-05	3148.0	Fri	0.0	11.16	0.0	15.30	1

Add time of year

Add time of year

```
In [81]: 1 df_bike['TimeOfYear'] = (df_bike.index - df_bike.index[0]).days / 365.0 # Days since the beginning of the year
        2 df_bike.head(3)
```

Out[81]:

	Total	DayOfWeek	IsHoliday	HoursDaylight	PRCP	TempC	IsDryDay	TimeOfYear
2012-10-03	3521.0	Wed	0.0	11.277359	0.0	13.35	1	0.000000
2012-10-04	3475.0	Thu	0.0	11.219142	0.0	13.60	1	0.002740
2012-10-05	3148.0	Fri	0.0	11.161038	0.0	15.30	1	0.005479

Generate and evaluate a model

Generate and evaluate a model

In [82]:

```
1 from sklearn.ensemble import GradientBoostingRegressor
2 from sklearn.dummy import DummyRegressor
3 from sklearn.metrics import mean_absolute_error
4
5 # drop any rows with missing data
6 df_bike.dropna(axis=0, how='any', inplace=True)
7
8 X_bike = pd.get_dummies(df_bike.loc[:,df_bike.columns != 'Total'])
9 display(X_bike.head(1).round(2))
10 y_bike = df_bike.Total
11
12 X_bike_train = X_bike.loc['2012':'2014']
13 y_bike_train = y_bike.loc['2012':'2014']
14 X_bike_test = X_bike.loc['2015']
15 y_bike_test = y_bike.loc['2015']
16
17 dummy_bike = DummyRegressor().fit(X_bike_train,y_bike_train)
18 gb_bike = GradientBoostingRegressor().fit(X_bike_train,y_bike_train)
19 print(f'dummy training mae      : {mean_absolute_error(y_bike_train,dummy_bike.predict(X_bike_train)).round(2)}')
20 print(f'one-back training mae : {mean_absolute_error(y_bike_train,y_bike_train.shift(1).fillna(0)).round(2)}')
21 print(f'gb training set mae     : {mean_absolute_error(y_bike_train,gb_bike.predict(X_bike_train)).round(2)}')
22 print(f'gb test set R^2         : {mean_absolute_error(y_bike_test,gb_bike.predict(X_bike_test)).round(2)}')
```


	IsHoliday	HoursDaylight	PRCP	TempC	IsDryDay	TimeOfYear	DayOfWeek_Fri	DayOfWeek_Mon	DayOfWeek_Sat	DayOfWeek_Sun	DayOfWeek_Thu	DayOfWeek_Tue	DayOfWe
2012-10-03	0.0	11.28	0.0	13.35	1	0.0	0	0	0	0	0	0	1

```
dummy training mae      : 1019.45
one-back training mae   : 710.39
gb training set mae     : 213.37
gb test set R^2         : 308.51
```

Plot predictions vs observed

Plot predictions vs observed

```
In [83]: 1 df_bike['predicted'] = gb_bike.predict(X_bike)
2 df_bike[['Total', 'predicted']].rolling(30, center=True, win_type='gaussian').mean(std=7).plot(alpha=0.5, figsize=(24, 8))
3 plt.tight_layout()
```



Time Series Operations Review

- Shifting
- Resampling
 - Downsampling
 - Upsampling
- Moving/Rolling Windows
- for more info, including time-series cross-validation:
 - [sklearn: Time-related feature engineering](#)
 - PML Chapter 13 - Modeling Sequential Data Using Recurrent Neural Network (with Tensorflow)
- for more models:
 - [skforecast](#)
 - [statsmodels](#)

Questions re Time Series Transformations?